

Parallel Algorithms

by

Biing-Feng Wang

王 炳 豐

Department of Computer Science
National Tsing Hua University

References:

- [1] J. JáJá, *An Introduction to Parallel Algorithms*, Addison Wesley, 1992.
- [2] S. G. Akl, *Parallel Computation: Models and Methods*, Prentice-Hall, 1997.
- [3] Journals and proceedings

An Example

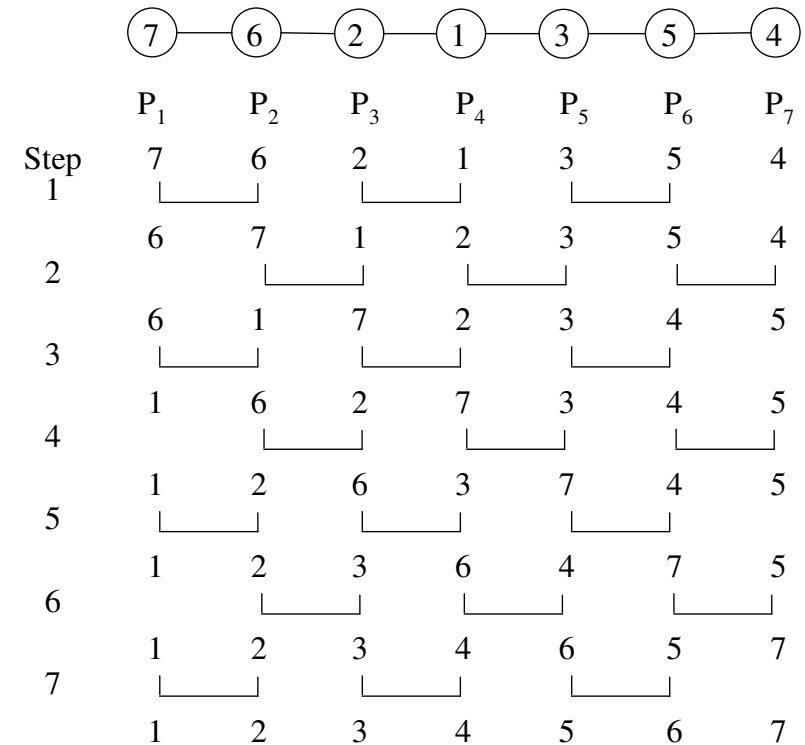
Problem: Sorting (Odd-Even Transposition Sort)

Input : 7, 6, 2, 1, 3, 5, 4

Output : 1, 2, 3, 4, 5, 6, 7

Model : Linear Processor Array ($n = 7$)

○ processor
— link



algo. { correctness ???
time complexity ???

Theorem: Odd-Even transposition sort produces a sorted sequence of n data after n steps.

$O(n)$ time ~~$\Omega(n \lg n)$~~ lower bound

Outline

Part I

1. Introduction
 - Classification of Parallel Computers
 - Performance of Parallel Algorithms
2. Shared-Memory Computers, Basic Techniques, and Brent's Theorem

Part II

3. Tree Machines
4. Linear Processor Arrays
5. Mesh-Connected Computers
6. Hypercubes
7. Perfect Shuffles
8. Mesh-Connected Computers with Multiple Broadcasting
9. Processor Arrays with Reconfigurable Bus Systems

Part III

10. Systolic Architectures
11. Randomized Algorithms and P-Completeness

Grading

- | | | |
|------------------------------|----------|---------------|
| • Midterm Examination | 50% | } ⇒ normalize |
| • Final Examination | 40% | |
| • Image | 10% | |
| • <u>Phone/Notebook</u> | -5%/each | |
| • <u>Fail if 3+ absences</u> | | |

Office Hours: 14:30 ~ 15:30, Monday to Friday
(R548, EECS)

Webpage: eLearn

Introduction


■ Classification of Parallel Computers

1. Flynn's Classification
(M. J. Flynn, "Very high speed computing systems,"
Proceedings of the IEEE, vol. 54, 1966, pp. 1901-1909)

Instruction stream (*I*): a sequence of instructions performed by a computer

Data stream (*D*): a sequence of data used to execute an instruction stream

SISD: single instruction stream, single data stream




One instruction is performed at a time, on one set of data.

practical,
easy to implement


INTRO-5

SIMD: single instruction stream, multiple data streams




One type of instruction is performed at a time, possible on different sets of data.

MISD: multiple instruction streams, single data stream



Different instructions on the same data can be performed at a time.

MIMD: multiple instruction streams, multiple data streams



Different instructions on different data can be performed at a time.

1-5e


1-5f

INTRO-6

2. Schwartz's Classification


(J. T. Schwartz, "Ultra-computers", *ACM Transactions on Programming Languages and Systems*, vol. 2, no. 4, 1980, pp. 484-521.)

1. *Paracomputer* (*shared-memory computer*)



- * Communication between any two processors takes $O(1)$ time through the shared memory
- * The **SIMD shared-memory** computer is also called *Parallel Random Access Machine (PRAM)*


2. *Ultracomputer*



- * The processors communicate with one another through an interconnection network. \Rightarrow one time unit (as fast as memory access)
 \neq distributed systems (computer networks) (distributed algorithms)

Examples of widely used interconnection network:


1. Linear Processor Array ($n = 7$)



diameter
maximum degree
link complexity


2. Mesh-Connected Computer (2d 3x3 mesh) * k-d mesh: $d_1 \times d_2 \times \dots \times d_k$

* 1-d mesh: LPA




most practical
1st popular

3. Tree Machine ($n = 2^l - 1 = 2^4 - 1, l = \lg(n+1)$)




level = 4
(height = 3)

4. Hypercube (k -dimension $\Rightarrow n = 2^k, k = \lg n$)




1-cube

2^1 nodes



2-cube

2^2 nodes




3-cube

2^3 nodes

most important one
(most famous, the best, 10^{3+})

1-8a

5. Cube-connected Cycle network (each node of a k -cube is replayed by a cycle of k processors)



k -d $\Rightarrow n = 2^k \times k$

e.g., $k = 3$

$2^3 \times 3$ nodes

diameter := $O(k \times (k/2)) = O(k^2)$
(= $2k - 2 + \lfloor k/2 \rfloor$ for $k \geq 4$)

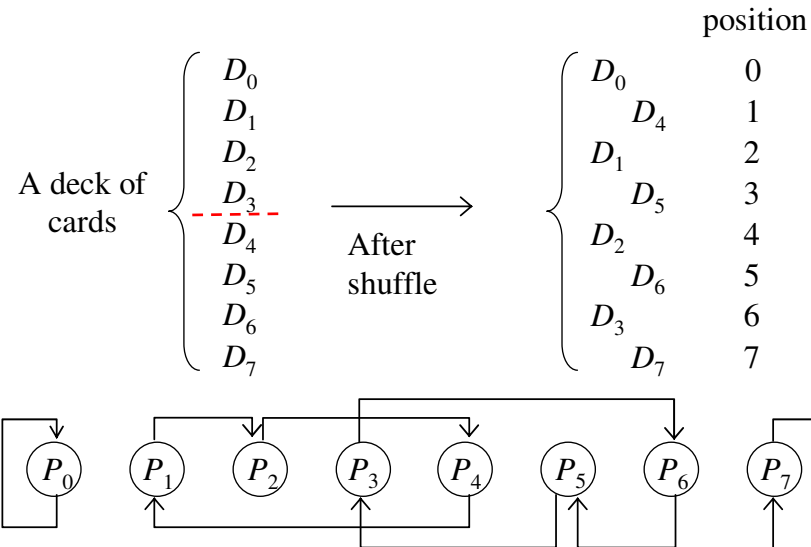


1-8b

1-8d


6. *Perfect Shuffle*

$$k-d \Rightarrow n = 2^k$$

most interesting one

7. *Shuffle-Exchange Network*8. *Single-channel broadcast communication System*

* Broadcast a data requires $O(1)$ time. pros: simple, $O(1)$ diameter
cons: one at a time

9. *Mesh-Connected Computer with Multiple Broadcasting*

* Current PC ???

1-10a

■ Performance of Parallel Algorithms

* Assume that you are familiar with O , θ , Ω

$$\text{Speedup} = \frac{\text{worst-case running time of fastest known sequential algorithm}}{\text{worst-case running time of parallel algorithm}}$$

1-11a

*A parallel algorithm is said to achieve *linear speedup* if the speedup with p processor is $\theta(p)$.

Question: super linear speedup ?

Cost = (parallel running time) × (number of processors used)

1-11a

* A parallel algorithm is said to be *cost-optimal* if the cost matches the (sequential) *lower bound* to within a constant multiplicative factor.

→ fastest known
目前已知最快
(是否 optimal 會隨時間改變)

* Which one, speedup or cost, is more important ???

worst-case running time of fastest known
sequential algorithm

Efficiency = $\frac{\text{worst-case running time of fastest known sequential algorithm}}{\text{cost of parallel algorithm}}$

1-11a

* used only for
experiments

cost of parallel algorithm

= (speedup) / (number of used processors)

≤ 1 (100%, used for comparing the real running time)

→ $\theta(1)$, an optimal algo has an efficiency of $\theta(1)$

Example: Consider the odd-even transposition sort on a linear array of n processors, which sort n data in $O(n)$ time.

1-11c

Speedup = $(n \log n) / n = \log n$ (all in θ -notation)

cost = $n \times n = n^2$ (non-optimal!)

efficiency = $\log n / n$

* Note: treat all O as θ

Another example:

sort n numbers in $O(n^{0.5})$ time using $O(n^{0.5} \lg n)$ PEs

Question 1: cost optimal → linear speedup ?

Question 2: cost optimal ← linear speedup ?

cost optimal ↔ linear speedup ↔ efficiency = $\theta(1)$